

6DoF Pose Estimation from Mesh using Differentiable Rendering

Meet Gera Abhinav Raundhal Amey Karan

International Institute of Information Technology, Hyderabad, India

December 7, 2024

Outline

- 1 Introduction
- 2 Differentiable Rendering Approach
- 3 Implementation Details
- 4 Inferences
- 5 Conclusion
- 6 Contributions
- 7 References

Introduction

- **Pose Estimation:**

- Determine position and orientation (6DoF) of objects.
- Crucial in robotics, augmented reality, computer vision.

- **Objective:**

- Optimize pose parameters using **Differentiable Rendering**.
- Compare different rotation representations:
 - Axis-Angle
 - Quaternions
 - Rotation Matrices

- **Approach:**

- Use silhouette loss between rendered and target images.
- Employ gradient-based optimization.

Differentiable Rendering

- **Definition:**

- Rendering process where gradients can be computed w.r.t input parameters.

- **Advantages:**

- Enables optimization of 3D scene parameters.
- Utilize standard backpropagation for parameter updates.

- **Applications:**

- Pose estimation
- Shape optimization
- Material property estimation

Problem Formulation

- **Given:**

- 3D mesh \mathcal{M} .
- Target image \mathcal{I}_{GT} .

- **Objective:**

- Estimate pose $\mathbf{T} = (\mathbf{R}, \mathbf{t})$ to align \mathcal{M} with \mathcal{I}_{GT} .

- **Rendering Function:**

$$\mathcal{I}_{\text{rendered}} = \mathcal{R}(\mathcal{M}, \mathbf{T})$$

- **Optimization Problem:**

$$\min_{\mathbf{T}} \mathcal{L}(\mathcal{I}_{\text{rendered}}, \mathcal{I}_{GT})$$

Silhouette Loss Function

- **Loss Function:**

$$\mathcal{L} = \sum_{i,j} \left(s_{i,j}^{\text{rendered}} - s_{i,j}^{\text{GT}} \right)^2$$

- **Where:**

- $s_{i,j}^{\text{rendered}}$: Pixel value at (i,j) in rendered silhouette.
- $s_{i,j}^{\text{GT}}$: Pixel value at (i,j) in ground truth silhouette.

- **Purpose:**

- Measures alignment between rendered and target silhouettes.

Gradient Computation

- **Goal:**

- Compute gradient of loss w.r.t pose parameters.

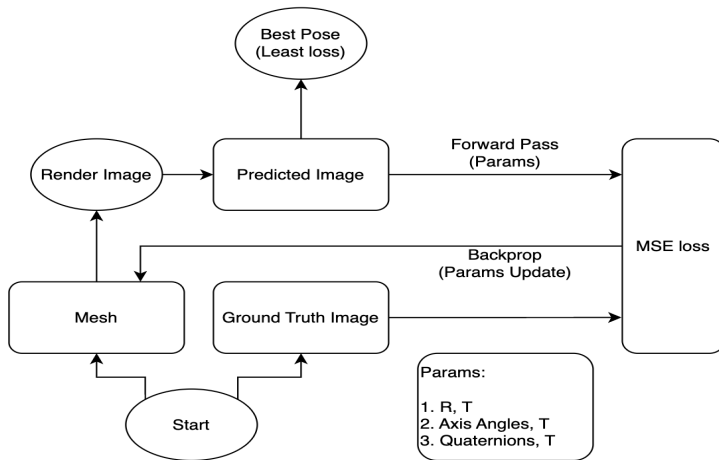
- **Gradient of Loss:**

$$\frac{\partial \mathcal{L}}{\partial q_i} = \frac{\partial \mathcal{L}}{\partial I_p} \cdot \frac{\partial I_p}{\partial R} \cdot \frac{\partial R}{\partial q_i}$$

- **Where:**

- q_i : Pose parameter (e.g., quaternion component).
- I_p : Projected image.
- R : Rotation matrix.

Optimization Process



Optimization Process

1 Initialization:

- Set initial estimates for \mathbf{R} and \mathbf{t} .
- Choose rotation representation (Axis angle, rotation matrix, quaternion).

2 Forward Pass:

- Render image I_p using current parameters.
- Compute loss \mathcal{L} .

3 Backward Pass:

- Compute gradients w.r.t parameters.

4 Parameter Update:

- Update parameters using optimizer (e.g., Adam).
- Apply constraints (normalization, orthogonality).

5 Iterate until Convergence.

Parameter Constraints

- **Quaternion Normalization:**

$$\mathbf{q} \leftarrow \frac{\mathbf{q}}{\|\mathbf{q}\| + \epsilon}$$

- **Rotation Matrix:** Should remain orthonormal at each step.
- **Purpose:**
 - Ensure parameters represent valid rotations.
 - Maintain numerical stability.

Avoiding Local Minima

- **Problem:**

- Optimization may get stuck in local minima.

- **Solution:**

- Introduce perturbations when progress stalls.

- **Methods:**

- **Axis-Angle:**

- Add small random vector to axis-angle parameters.

- **Quaternions:**

- Perturb and re-normalize quaternion.

- **Rotation Matrices:**

- Convert to axis-angle, perturb, convert back.

Results and Observations

Results can be viewed here: [Github](#)

Results and Observations

- **Loss Convergence:**

- All models showed convergence.
- Rotation matrices had more fluctuations.

- **Parameter Updates:**

- Translation and rotation parameters converged exponentially.
- Perturbations caused spikes in updates.

- **Comparison:**

- Axis-Angle and Quaternions converged faster.
- Fewer parameters lead to quicker convergence.

Interpretation of Results

- **Number of Parameters:**

- Axis-Angle: 3+3 parameters.
- Quaternions: 4+3 parameters (normalized).
- Rotation Matrices: 9+3 parameters (with constraints).

- **Impact on Optimization:**

- Higher dimensionality increases computational load.
- Constraints add complexity.

- **Recommendation:**

- Use Axis-Angle or Quaternions for efficiency.

Conclusion

- **Differentiable Rendering:**

- Effective for refining pose estimation.
- Enables gradient-based optimization of pose parameters.

- **Rotation Representations:**

- Choice impacts convergence and stability.
- Axis-Angle and Quaternions preferred for efficiency.

- **Future Work:**

- Explore hybrid approaches.
- Investigate other loss functions and regularizations.

Contribution of Team Members

• **Meet Gera:**

- Implemented Differential Rendering Algorithm.
- Derived mathematical expressions for backpropagation.
- Contributed to report and documentation.

• **Abhinav Raundhal:**

- Implemented Differential Rendering Algorithm.
- Conducted experiments on different objects.
- Contributed to report and documentation.

• **Amey Karan:**

- Set up environment and dependencies.
- Implemented Differential Rendering Algorithm.
- Contributed to report and documentation.

References I



I. Shugurov, I. Pavlov, S. Zakharov, and S. Ilic, “Multi-view object pose refinement with differentiable renderer,” *IEEE Robotics and Automation Letters*, vol. 6, pp. 2579–2586, Apr. 2021.



A. Palazzi, L. Bergamini, S. Calderara, and R. Cucchiara, “End-to-end 6-dof object pose estimation through differentiable rasterization,” in *Computer Vision – ECCV 2018 Workshops* (L. Leal-Taixé and S. Roth, eds.), (Cham), pp. 702–715, Springer International Publishing, 2019.



N. Ravi, J. Reizenstein, D. Novotny, T. Gordon, W.-Y. Lo, J. Johnson, N. Neverova, G. Larsson, R. Girdhar, M. Singh, S. Somasundaram, A. Makadia, and M. Rohrbach, “Pytorch3d: An optimized library for 3d deep learning.”

<https://github.com/facebookresearch/pytorch3d>, 2020.

GitHub repository.